




Pre-trained CNNs

186


Pre-trained CNNs

Pre-trained CNNs

- Pre-trained architectures
 - VGG
 - RESNET
 - Inception
 - DenseNet
 - EfficientNet

187

Pre-trained CNNs




- Pre-trained CNNs refer to CNN models that have been trained on large datasets, and then made available for transfer learning or fine-tuning on other tasks.
- These pre-trained CNNs serve as feature extractors, capturing hierarchical representations of features from raw input data, which can then be used for a variety of downstream tasks such as image classification, object detection, and image segmentation.
- These pre-trained CNN models are typically available in deep learning frameworks such as TensorFlow, PyTorch, and Keras.
- By using pre-trained CNNs, developers can benefit from the features learned by these models on large datasets and fine-tune them on smaller, task-specific datasets to achieve better performance with less computational resources and training time.

188

Pre-trained CNNs

The most common Pre-trained CNNs:



- **VGG (Visual Geometry Group)**
- **ResNet (Residual Network)**
- **Inception (GoogLeNet)**
- **DenseNet**
- **EfficientNet**


189

Pre-trained CNN architectures

VGG (Visual Geometry Group)

VGG networks are characterized by their simplicity and uniform architecture, consisting of multiple convolutional layers followed by max-pooling layers and fully connected layers.

Variants of VGG, such as VGG16 and VGG19, are often used as feature extractors.



190

VGG

Advantages

Simple and Uniform Architecture: VGG networks have a straightforward and uniform architecture, consisting of stacked convolutional layers followed by max-pooling layers and fully connected layers. Easy to understand, implement, and modify.

State-of-the-Art Performance: Its deep architecture allowed it to learn rich hierarchical representations of visual features, leading to high accuracy in recognizing objects within images.

Transfer Learning: Pre-trained VGG models, trained on large-scale datasets such as ImageNet, are widely available and commonly used as feature extractors or as the basis for transfer learning in various computer vision tasks.

Effective Feature Learning: With their deep architecture and small 3×3 convolutional filters, excel at learning complex and abstract features from input images.

Generalization: VGG networks generalize well to unseen data, thanks to their depth and hierarchical feature learning capabilities.

191

VGG

Disadvantages

High Computational Cost: The large number of parameters and operations in the network require significant computational resources, memory, and time, making them less suitable for resource-constrained environments.

Memory Requirements: The deep architecture of VGG networks results in high memory requirements, both during training and inference.

Vanishing Gradient Problem: Deeper VGG variants, such as VGG19, suffer from the vanishing gradient problem during training, where gradients become extremely small or vanish as they propagate backward through the network.

Overfitting: VGG networks, especially the deeper variants, are prone to overfitting when trained on small datasets or when insufficient regularization techniques are applied.

Lack of Spatial Attention: VGG networks do not explicitly model spatial attention mechanisms, which can limit their ability to focus on relevant regions of an image while ignoring irrelevant information.


192

Pre-trained CNN architectures

ResNet (Residual Network)

ResNet introduced the concept of residual learning, which helps in training very deep networks by mitigating the vanishing gradient problem.

ResNet architectures, such as ResNet50, ResNet101, and ResNet152, have been pre-trained on ImageNet and are widely used for transfer learning.



193

ResNet

Advantages

Addressing Vanishing Gradient Problem: Residual connections, introduced in ResNets, mitigate the vanishing gradient problem encountered in very deep neural networks.

Ease of Training: ResNets are relatively easier to train compared to earlier deep architectures, such as VGG or plain deep networks.

Improved Performance: Their ability to capture rich hierarchical representations of features leads to higher accuracy and better generalization on challenging datasets like ImageNet.

Scalability: ResNets are highly scalable architectures, capable of accommodating a wide range of depths.

Transfer Learning: Pre-trained ResNet models, trained on large-scale datasets like ImageNet, are readily available and widely used as feature extractors or as the basis for transfer learning in various computer vision tasks.

194

ResNet

Disadvantages

Increased Computational Complexity: Although ResNets address the vanishing gradient problem and enable training of deeper networks, they come with increased computational complexity compared to shallower architectures.


Memory Consumption: Deeper ResNet variants require more memory to store activations and gradients during training.

Overfitting: While ResNets are generally robust to overfitting, deeper variants may still be prone to overfitting, particularly when trained on small datasets or with insufficient regularization. Techniques such as dropout, weight decay, and data augmentation may be necessary to mitigate overfitting effectively.

Design Complexity: Designing and fine-tuning ResNet architectures require careful consideration of hyperparameters such as the number of layers, the size of convolutional filters, and the learning rate schedule.

Interpretability: Deeper ResNet variants with numerous layers may be less interpretable compared to shallower architectures.

195



Pre-trained CNN architectures

Inception (GoogLeNet)

The Inception architecture uses multiple parallel convolutional paths with different filter sizes to capture features at different scales.

Models like InceptionV3 and InceptionResNetV2 are popular pre-trained CNNs that have been trained on large datasets.

196

Inception Advantages

Efficient Use of Parameters: Allows for the efficient use of parameters by performing convolutions with multiple filter sizes (1×1 , 3×3 , and 5×5) in parallel.

Increased Computational Efficiency: Achieve high performance with fewer parameters and computations compared to other architectures like VGG or ResNet.

Improved Performance: Inception networks have achieved state-of-the-art performance on various computer vision tasks, including image classification, object detection, and segmentation. Their ability to capture multi-scale features and hierarchical representations leads to higher accuracy and better generalization on challenging datasets.

Scalability: Are highly scalable architectures, capable of accommodating different depths and widths to balance computational efficiency and performance.

Interpretability: Inception networks produce highly interpretable feature maps due to the use of parallel convolutional operations at different scales.

197

Inception Disadvantages

Complexity: The inception module introduces additional complexity to the architecture, making it more challenging to design and optimize compared to simpler architectures like VGG or ResNet.

Memory Consumption: Inception networks may require more memory to store activations and gradients during training, especially when dealing with deep architectures and large datasets.

Overfitting: While Inception networks are generally robust to overfitting, deeper variants may still be prone to overfitting, particularly when trained on small datasets or with insufficient regularization.

Transfer Learning Complexity: While pre-trained Inception models are available for transfer learning, fine-tuning Inception networks may require additional effort and experimentation due to the complexity of the architecture.

Deployment Challenges: Deploying Inception networks on embedded or mobile devices may pose challenges due to their computational and memory requirements.


198

Pre-trained CNN architectures

DenseNet

DenseNet connects each layer to every other layer in a feed-forward fashion, which helps in feature reuse and encourages feature propagation.

DenseNet models, such as DenseNet121 and DenseNet169, are commonly used for transfer learning tasks.



199

DenseNet

Advantages

Feature Reuse: Facilitates extensive feature reuse by connecting each layer to every other layer in a feed-forward fashion within each dense block. This dense connectivity pattern allows for better information flow through the network, enabling the efficient reuse of features learned at different depths.

Addressing Vanishing Gradient: Mitigates the vanishing gradient problem by providing direct paths for gradient flow from the final layer back to earlier layers through dense connections.

Parameter Efficiency: Achieves high model capacity and expressiveness with fewer parameters compared to traditional architectures.

Improved Gradient Flow: The dense connectivity pattern allows gradients to flow directly through the network, facilitating smoother and more efficient optimization during training.

Feature Diversity: Encourages the learning of diverse and complementary features at different depths of the network due to the dense connections.

200

DenseNet

Disadvantages

Computational Complexity: Architectures can be computationally expensive, especially when dealing with deep networks and large datasets.

Memory Consumption: Architectures may require more memory to store activations and gradients during training, particularly when dealing with deep architectures and large batch sizes.

Deployment Challenges: Deploying models on resource-constrained devices or platforms may pose challenges due to their computational and memory requirements.


Design Complexity: Designing and fine-tuning DenseNet architectures require careful consideration of hyperparameters such as the number of layers, the growth rate, and the size of dense blocks. Optimizing these hyperparameters can be challenging and may require extensive experimentation.

Interpretability: Architectures with dense connectivity patterns may produce less interpretable feature maps compared to architectures with more structured connections. Understanding the learned representations and decision-making processes in DenseNet models can be challenging.

201

Pre-trained CNN architectures

EfficientNet



EfficientNet is a family of models that achieve state-of-the-art accuracy with significantly fewer parameters and FLOPs (floating-point operations) compared to other architectures.

EfficientNet models are efficient in terms of both accuracy and computational cost and are widely used for transfer learning.

202

EfficientNet

Advantages

Efficient Parameter Usage: Achieves high model capacity and expressiveness with fewer parameters compared to traditional architectures like VGG, ResNet, or DenseNet.

Improved Computational Efficiency: Architectures are computationally efficient, making them suitable for deployment on resource-constrained devices and platforms.

Scalability: Architectures are highly scalable, capable of accommodating different depths, widths, and resolutions to balance computational efficiency and performance.

State-of-the-Art Performance: Achieves state-of-the-art performance on various computer vision tasks, including image classification, object detection, and segmentation. Its ability to capture rich hierarchical representations of features leads to higher accuracy and better generalization on challenging datasets like ImageNet.

Transfer Learning: Pre-trained EfficientNet models, trained on large-scale datasets like ImageNet, are readily available and widely used as feature extractors or as the basis for transfer learning in various computer vision tasks. Transfer learning with EfficientNet can significantly reduce the time and resources required to train models from scratch.

203

EfficientNet

Disadvantages

Complexity: Architectures are more complex compared to simpler architectures like VGG or ResNet. Designing and fine-tuning EfficientNet architectures require careful consideration of hyperparameters such as depth, width, resolution, and scaling coefficients.

Memory Consumption: While EfficientNet architectures are more memory-efficient compared to traditional architectures, they may still require significant memory to store activations and gradients during training, especially when dealing with deep architectures and large batch sizes.

Overfitting: While EfficientNet architectures are generally robust to overfitting, deeper or larger variants may still be prone to overfitting, particularly when trained on small datasets or with insufficient regularization.

Deployment Challenges: Deploying models on embedded or mobile devices may pose challenges due to their computational and memory requirements.

Interpretability: Architectures may produce less interpretable feature maps compared to architectures with more structured connections. Understanding the learned representations and decision-making processes in EfficientNet models can be challenging.

204

Embedding Learning

205

Embedding Learning

What is Embedding Learning? What is Embedding's?

- Why are embeddings used for ?
- How to Embedding?
- What can be embedded?
- How embeddings are created
- Types of Embeddings
- Real world Applications

206

Embedding Learning

Embedding is a critical tool for ML engineers who build text and image search engines, recommendation systems, chatbots, fraud detection systems and many other applications. In essence, embedding enables machine learning models to find similar objects.

Unlike other ML, embeddings are learned from data using various algorithms, such as neural networks, instead of explicitly requiring human expertise to define. They allow the model to learn complex patterns and relationships in the data, which would otherwise be impossible for humans to identify.

207