

ADVANCED TOPICS IN DATABASES



MongoDB Commands and Queries

Master in Informatics Engineering
Data Engineering

Informatics Engineering Department

MongoDB: Terminology

| RDBMS | MongoDB |
|-------------------|------------------|
| database instance | MongoDB instance |
| schema | database |
| table | collection |
| row | document |
| rowid | |

```
{
  "nome": "Sérgio Conceição",
  "nridentificacaocivil": 937587,
  "nif": 104052455,
  "datanascimento": "52.10.06"
},
{
  "nome": "Fernando Santos",
  "nridentificacaocivil": 937587,
  "nif": 104052455,
  "datanascimento": "52.10.06"
},
{
  "nome": "António Oliveira",
  "nridentificacaocivil": 937587,
  "nif": 104052455,
  "datanascimento": "52.10.06"
}
```

collection

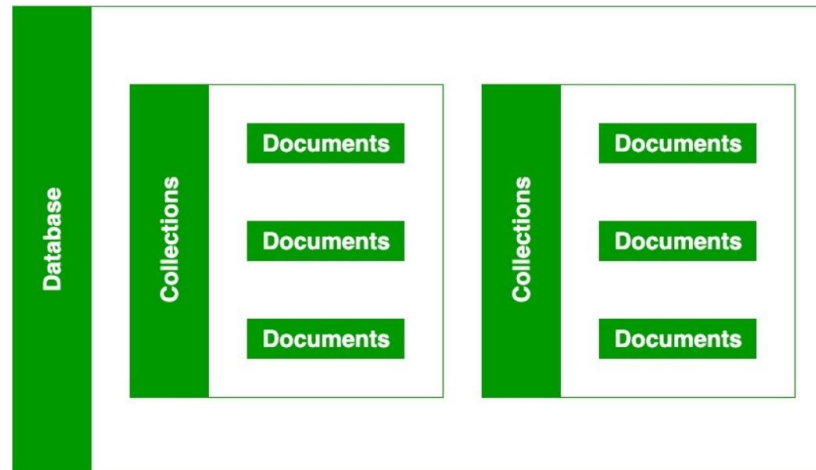
- each JSON **document**:
 - belongs to a **collection**
 - has a field **_id**
 - unique within the collection
- each collection:
 - belongs to a “**database**”

- <http://www.mongodb.org/>



Databases e Collections

- Each instance of MongoDB can manage multiple databases
- Each database is composed of a set of collections
- Each collection contains a set of documents
 - The documents of each collection represent similar“objects”
 - can enforce document validation rules for a collection during update and insert operations.



Databases

- Show the list of available databases

```
show databases
```

- Select the database you are interested in

```
use <database-name>
```

E.g. --> use MySchemadb

- Create a database and a collection inside the database
 - To select the database by using the command “use <database name>”
 - then, we can create a collection
 - MongoDB creates a collection implicitly when the collection is first referenced in a command



Databases

- Delete/Drop a database
 - Select the database by using “use <database name>”
 - Execute the command

E.g.,

```
db.dropDatabase()
```

```
use MySchemadb;
```

```
db.dropDatabase();
```



Collections

- A collection stores documents, uniquely identified by a document “_id”
- Create collections

```
db.createCollection(<collection name>, <options>);
```

- The collection is associated with the current database. Always select the database before creating a collection.
- Options related to the collection size and indexing, e.g., to create a capped collection, or to create a new collection that uses document validation

```
db.createCollection("authors", {capped: true});
```

```
db.createCollection( "logs", { capped: true, size: 500000 } ); // size is in bytes.
```

- Show collections

```
show collections
```

- Drop collections

```
db.<collection_name>.drop()
```

```
db.authors.drop()
```



C.R.U.D. Operations

• Create

```
db.users.insertOne(
  {
    name: "sue",
    age: 26,
    status: "pending"
  }
)
```

collection
field: value
field: value
field: value } document

• Read

```
db.users.find(
  { age: { $gt: 18 } },
  { name: 1, address: 1 }
).limit(5)
```

collection
query criteria
projection
cursor modifier

• Update

```
db.users.updateMany(
  { age: { $lt: 18 } },
  { $set: { status: "reject" } }
)
```

collection
update filter
update action

• Delete

```
db.users.deleteMany(
  { status: "reject" }
)
```

collection
delete filter



Inserts

```
db.inventory.insert( { _id: 10, type: "misc", item: "card", qty: 15 } )
```

Inserts a document with three fields into collection **inventory**

User-specified **_id** field

```
db.inventory.insert( { type: "book", item: "journal" } )
```

The database generates **_id** field

```
db.inventory.find()
```

```
{ "_id": ObjectId("58e209ecb3e168f1d3915300"), type: "book", item: "journal" }
```



Inserts

```
db.<collection name>.insertOne( {<set of the field:value pairs of the new document>} );
```

Example1

```
db.people.insertOne( {  
  user_id: "abc124",  
  age: 45,  
  favorite_colors: ["blue", "green"]  
} );
```

Favorite_colors is an array

Example2

```
db.people.insertOne( {  
  user_id: "abc124",  
  age: 45,  
  address: {  
    street: "my street",  
    city: "my city"  
  }  
} );
```

Nested document

```
db.<collection name>.insertMany([ <comma separated list of documents> ]);
```

```
db.products.insertMany( [  
  { user_id: "abc123", age: 30, status: "A"},  
  { user_id: "abc456", age: 40, status: "A"},  
  { user_id: "abc789", age: 50, status: "B"}  
] );
```



Delete

- Delete existing data, in MongoDB corresponds to the deletion of the associated document.
 - Conditional delete
 - Multiple delete

| | |
|--|--|
| <code>DELETE FROM people WHERE status = "D"</code> | <code>db.people.deleteMany({ status: "D" })</code> |
| <code>DELETE FROM people</code> | <code>db.people.deleteMany({})</code> |



Updates

- E.g.,

```
db.inventory.updateMany(  
  { "qty": { $lt: 50 } },  
  {  
    $set: { "size.uom": "in", status: "P" },  
    $currentDate: { lastModified: true }  
  }  
)
```

This operation updates all documents with qty <50

It sets the value of the size.uom field to "in", the value of the status field to "P", and the value of the lastModified field to the current date.



Updates

```
db.inventory.update(  
    { type: "book", item : "journal" },  
    { $set: { qty: 10 } },  
    { upsert: true } )
```

Finds all docs matching query

```
{ type: "book", item : "journal" }
```

and sets the field { qty: 10 }

upsert: true

if no document in the **inventory** collection matches

creates a new document (generated **_id**)

it contains fields **_id**, **type**, **item**, **qty**

MongoDB operator

```
db.<table>.updateMany(  
    { <condition> },  
    { $set: {<statement>} }  
)
```



Querying: Basics

- Mongo query language
- A MongoDB **query**:
 - Targets a specific **collection** of documents
 - Specifies **criteria** that identify the returned documents
 - May include a **projection** to **specify** returned **fields**
 - May impose limits, sort, orders, ...
- Basic query - all documents in the collection:

| Oracle clause | MongoDB operator |
|---------------|------------------|
| SELECT | find() |

| | |
|-------------------------------|--------------------------|
| SELECT * FROM users | db.users. find () |
|-------------------------------|--------------------------|

```
db.<collection name>.find( {<conditions>}, {<fields of interest>} );
```



Querying: Example

Collection Query Criteria Modifier
`db.users.find({ age: { $gt: 18 } }).sort({age: 1 })`

| |
|-----------------|
| { age: 18, ...} |
| { age: 28, ...} |
| { age: 21, ...} |
| { age: 38, ...} |
| { age: 18, ...} |
| { age: 38, ...} |
| { age: 31, ...} |

users

Query Criteria

| |
|-----------------|
| { age: 28, ...} |
| { age: 21, ...} |
| { age: 38, ...} |
| { age: 38, ...} |
| { age: 31, ...} |

Modifier

| |
|-----------------|
| { age: 21, ...} |
| { age: 28, ...} |
| { age: 31, ...} |
| { age: 38, ...} |
| { age: 38, ...} |

Results

`db.collection.find()` gives back a cursor. It can be used to iterate over the result or as input for next operations.

- `cursor.sort()`
- `cursor.count()`
- `cursor.limit()`
- `cursor.max()`
- `cursor.min()`
- `cursor.pretty()`



Querying

| MySQL clause | MongoDB operator |
|--------------|---------------------------|
| COUNT | count() or find().count() |

| | |
|---|---|
| SELECT COUNT(*) FROM people | db.people.count() or db.people.find().count() |
| SELECT COUNT(*) WHERE status = "A" FROM people | db.people.count(status: "A") or db.people.find({status: "A"}).count() |
| SELECT COUNT(*) FROM people WHERE age > 30 | db.people.count({ age: { \$gt: 30 } }) |



Querying: Selection

```
db.<collection name>.find( {<conditions>}, {<fields of interest>} );
```

```
db.inventory.find({ type: "snacks" })
```

All documents from collection **inventory** where the **type** field has the value **snacks**

```
db.inventory.find({ type: { $in: [ 'food', 'snacks' ] } })
```

All **inventory** docs where the **type** field is either **food** or **snacks**

```
db.inventory.find( { type: 'food', price: { $lt: 9.95 } } )
```

All ... where the **type** field is **food** and the **price** is **less than 9.95**

| Name | Description |
|------------------|--|
| \$eq or : | Matches values that are equal to a specified value |
| \$gt | Matches values that are greater than a specified value |
| \$gte | Matches values that are greater than or equal to a specified value |
| \$in | Matches any of the values specified in an array |
| \$lt | Matches values that are less than a specified value |
| \$lte | Matches values that are less than or equal to a specified value |
| \$ne | Matches all values that are not equal to a specified value, including documents that do not contain the field. |
| \$nin | Matches none of the values specified in an array |



Querying: Selection

| | |
|---|--|
| <pre>SELECT * FROM people WHERE age > 25</pre> | <pre>db.people.find({ age: { \$gt: 25 } })</pre> |
|---|--|

| | |
|--|--|
| <pre>SELECT * FROM people WHERE status = "A" OR age = 50</pre> | <pre>db.people.find({ \$or: [{ status: "A" } , { age: 50 }] })</pre> |
|--|--|

| | |
|--|--------------------|
| <pre>db.inventory.find({ item: null })</pre> | // equality filter |
|--|--------------------|

| | |
|--|---------------------|
| <pre>db.inventory.find({ item : { \$exists: false } })</pre> | // existence filter |
|--|---------------------|

| | |
|---|----------------|
| <pre>db.inventory.find({ item : { \$type: 10 } })</pre> | // type filter |
|---|----------------|

| | |
|---|---|
| <pre>SELECT * FROM people WHERE status = "A" AND age = 50</pre> | <pre>db.people.find({ status: "A", age: 50 })</pre> |
|---|---|

| | |
|--|---|
| | <pre>db.people.find({ "address.city": "Rome" })</pre> |
|--|---|

```
{ _id: "A",  
  address: {  
    street: "Via Torino",  
    number: "123/B",  
    city: "Rome",  
    code: "00184"  
  }  
}
```

nested document

Note:

- Item: null → matches documents that either contain the item field whose value is null or that do not contain the item field
- Item: { \$exists: false } → matches documents that do not contain the item field



Querying- examples

- Count

```
db.people. count({ age: 32 })
```

- Comparison

```
db.people. find({ age: { $gt: 32 } }) // or equivalently with $gte, $lt, $lte,
```

```
db.people.find({ age: { $in: [32, 40] } }) // returns all documents having age either 32 or 40
```

```
db.people.find({ age: { $gt: 25, $lte: 50 } }) //returns all documents having age > 25 and age <= 50
```

- Logical

```
db.people.find({ name: { $not: { $eq: "Max" } } })
```

```
db.people.find({ $or: [ {age: 32}, {age: 33} ] } )
```



Querying - examples

```
db.items.find({  
  $and: [  
    {$or: [{qty: {$lt: 15}}, {qty: {$gt: 50}} ]},  
    {$or: [{sale: true}, {price: {$lt: 5}} ]}  
  ]  
})
```

This query returns documents (items) that satisfy **both** these conditions:

1. Quantity sold either less than 15 **or** greater than 50
2. Either the item is on sale (field “sale”: true) **or** its price is less than 5



Querying - examples

- Embedded Documents

Select all documents where the field size equals the **exact document** { h: 14, w: 21, uom: "cm" }

```
db.inventory.find( { size: { h: 14, w: 21, uom: "cm" } } )
```

To specify a query condition on fields in an embedded/nested document, use **dot notation**

```
db.inventory.find( { "size.uom": "in" } )
```

Dot notation and comparison operator

```
db.inventory.find( { "size.h": { $lt: 15 } } )
```



Querying - examples

Array

Query for all documents where the field tags value is an array with exactly two specific elements

`db.inventory.find({ tags: ["red", "black"]})` → Item list **order is important**

`db.inventory.find({ tags: { $all: ["red", "black"]} })` → List order does **not** important

The following queries **return different** results, i.e., **they are not equivalent**

`db.inventory.find({ tags: ["red", "black"]})`

`db.inventory.find({ tags: ["black", "red"]})`

The following queries **return the same results**, i.e., **they are equivalent**

`db.inventory.find({ tags: { $all: ["red", "black"]} })`

`db.inventory.find({ tags: { $all: ["black", "red"]} })`



Aggregation Framework

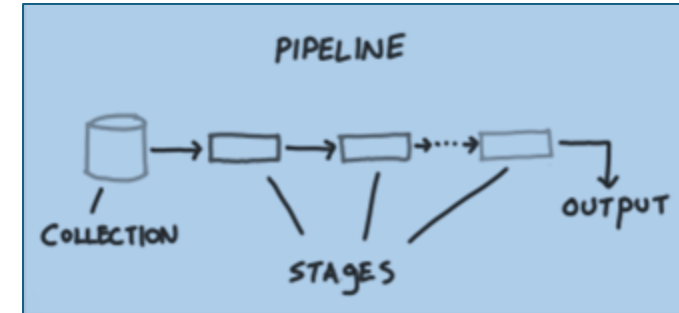
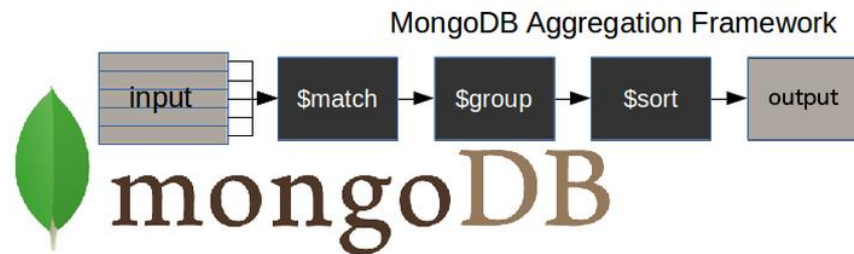
- Allows you to define data processing pipeline
 - Pipelines are made up of processing stages
 - Each step transforms the data as it moves through it
 - Each stage is a MongoDB aggregation framework command
 - Find documents
 - Modify the Structure of the Query Response Documents
 - Generate aggregated results
 - Modify Documents in DB
 - create new documents in the DB.

<https://docs.mongodb.com/manual/aggregation/#aggregation-framework>



Aggregation Framework

- A typical MongoDB aggregation pipeline.



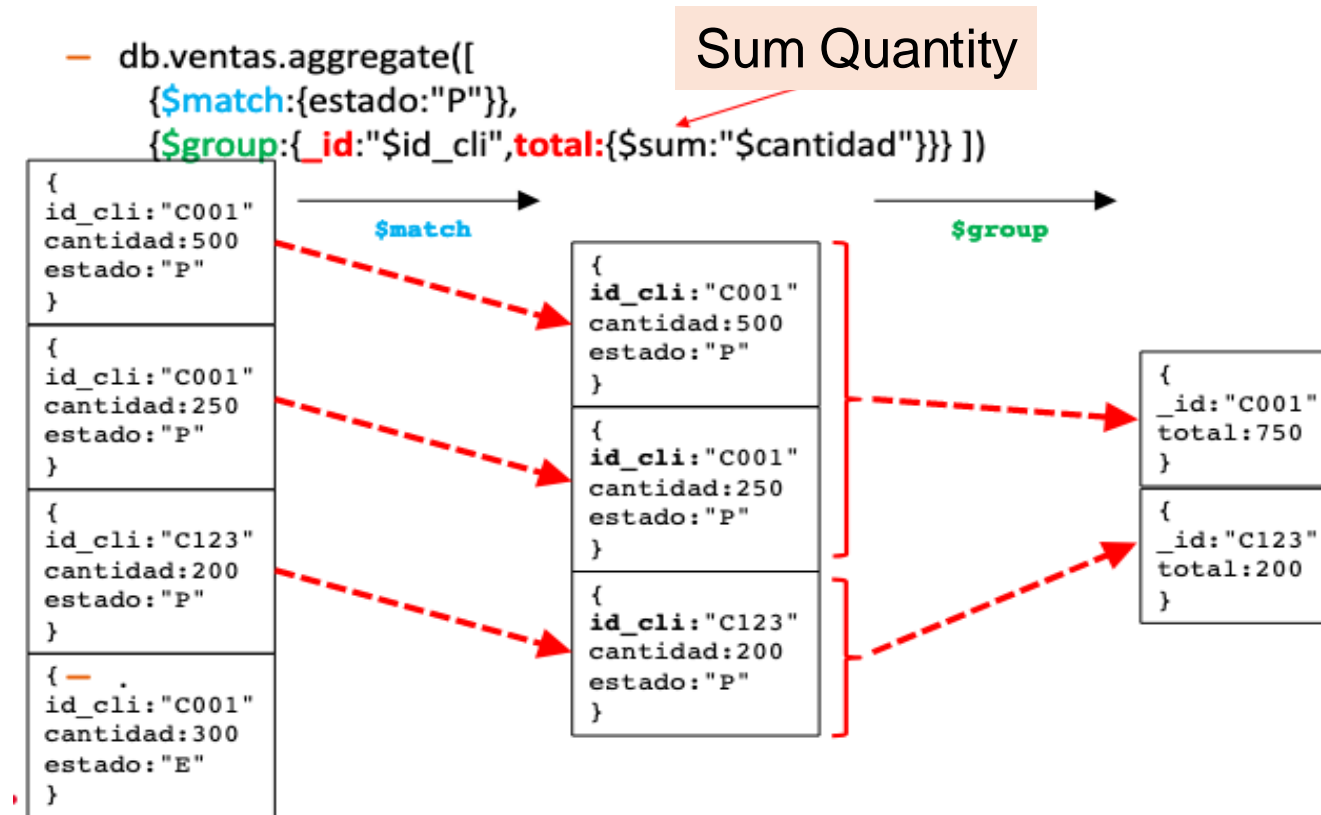
- \$match – filters the documents we need to work with, the ones that meet our needs
- \$group – does the aggregation work
- \$sort – classifies the resulting documents in the way we require (ascending or descending)

The pipeline input can be only one collection, where others can later be merged into the pipeline. The pipeline performs successive transformations in the data until our goal is achieved.



Aggregation Framework

Exemplification



Aggregation Framework

- General shape of the pipeline:

```
db.collection.aggregate( [ { <etapa1> }, { <etapa2> }, ... ] )
```

- Examples of types of steps that can be used with aggregate:

- \$match: Moves the documents in the stream that match the filter to the next stage of the pipeline.
- \$group: Group documents and applies aggregation operations on the groups;
- \$project: Allows you to add new attributes or remove existing attributes. Generates an output document for each input
- \$sort: sort the documents by the chosen key;
- \$limit: passes the first n input documents to the pipeline (where the chosen threshold is);
- \$unwind : Expands an array, generating an output document for each input in the array.



Groupings & Aggregations

Exemplo:

```
db.universities.aggregate([
  { $match : { name : 'Isep' } },
  { $unwind : '$students' },
  { $project : { _id : 0, 'students.year' : 1, 'students.number' : 1 } },
  { $sort : { 'students.number' : -1 } }
]).pretty()
```



Groupings & Aggregations

- A \$group step allows you to do groupings and aggregations

Examples:

- Sum of all product values in an order;
- Average grades of students in a subject.
- These operations resemble the group by, having, count, avg, sum, max, and min of the SQL language

<https://docs.mongodb.com/manual/aggregation/#aggregation-framework>



Simple aggregations

General form of simple aggregations (without grouping):

```
db.collection.aggregate(  
  { $group :  
    { _id : null,  
      <nome_resultado_agregação1>:{  
        <função_agregação1>:"$<atributo_a_agregar>"},  
      <nome_resultado_agregação2>:{  
        <função_agregação2>:"$<atributo_a_agregar>"},  
      ...  
    }  
  }  
);
```



Aggregation example

```
{
  "_id": "10280",
  "city": "NEW YORK",
  "state": "NY",
  "pop": 5574,
  "loc": [
    -74.016323,
    40.710537
  ]
}
```

```
db.zipcodes.aggregate( [
  { $group: { _id: null,
              totalPop: { $sum: "$pop" }
            }
  }
])
```

Query Result

```
{ "_id" : null, "totalPop" : 248408400 }
```



Agregations

General form of aggregations with clustering:

```
db.collection.aggregate(  
  { $group :  
    { _id : <atributo(s)_agrupador(es)>,  
      <nome_resultado_agregação1>:{  
        <função_agregação1>:"$<atributo_a_agregar>"},  
      <nome_resultado_agregação2>:{  
        <função_agregação2>:"$<atributo_a_agregar>"},  
      ...  
    }  
  }  
);
```



Example

- For each state, calculate the total population and the average population;

```
db.zipcodes.aggregate( [  
  { $group: { _id: "$state",  
              totalPop: { $sum: "$pop" },  
              mediaPop: { $avg: "$pop" }  
            }  
  }  
)
```



Example

- For each city in the state of New York, calculate the total population

```
db.zipcodes.aggregate(  
  [  
    { $match: { state: "NY"} },  
    { $group: { _id: "$city",  
                cityTotalPop: { $sum: "$pop" }  
            }  
    }  
  ]  
);
```

```
SELECT city, SUM(pop) AS cityTotalPop  
FROM zipcodes  
Where state like 'NY'  
GROUP BY city;
```



Example

```
db.zipcodes.aggregate( [  
  { $group: { _id: "$state",  
              totalPop: { $sum: "$pop" }  
            }  
  },  
  { $match: { totalPop: { $gte: 10*1000*1000 } } }  
)
```

```
SELECT state, SUM(pop) AS totalPop  
FROM zipcodes  
GROUP BY state  
HAVING totalPop >= (10*1000*1000)
```



Example

- Group by state and city

```
db.zipcodes.aggregate( [  
  { $group: { _id: {gstate:"$state",  
gcity:"$city"},  
              cityPop: { $sum: "$pop" }  
            }  
  }  
])
```



Example

1. List user names in uppercase letters and in alphabetical order.

```
db.users.aggregate(  
  [  
    { $project : { name:{$toUpper:"$_id"} , _id:0 } },  
    { $sort : { name : 1 } }  
  ]  
)
```

```
{  
  "name" : "JANE"  
},  
{  
  "name" : "JILL"  
},  
{  
  "name" : "JOE"  
}
```

The operator - \$project:

- Creates a new field called Name.
- Converts the value from _id to uppercase, with the \$toUpper operator. Then \$project creates a new field, called name to hold this value.
- Suppresses the ID field. The \$project will pass the _id field by default unless explicitly suppressed. •
- The \$sort operator sorts the results by the name field.



Example

2. List how many people have joined the club in each month of the year.

```
db.users.aggregate(  
  [  
    { $project : { month_joined : { $month : "$joined" } } },  
    { $group : { _id : {month_joined:"$month_joined"} , number : { $sum : 1 } } },  
    { $sort : { "_id.month_joined" : 1 } }  
  ]  
)
```

- The **\$project** operator creates a new attribute called month_joined.

The **\$month** operator converts the attribute values to integer representations of the month (1,2,...). The \$project operator then assigns the values to **the month_joined attribute**.

The operator **\$group** collects all documents with a certain value of month_joined and counts how many documents there are for that value. Specifically, for each unique value, **\$group creates a new document "per month"** with **two attributes**:

- **_id**, which contains a document with the month_joined attribute and its value.

- **number**, which is a generated attribute. **The \$sum operator increments this attribute by 1 for each document containing the value month_joined provided.**



Example of a two-step grouping & aggregation

```
db.zipcodes.aggregate( [  
  { $group: { _id: { state: "$state", city: "$city" },  
               cityPop: { $sum: "$pop" }  
            }  
  },  
  { $group: { _id: "$_id.state",  
               avgCityPop: { $avg: "$cityPop" }  
            }  
  }  
])
```

This block groups zip codes by state and city and calculates the total population of each state.

Block result: set of triples (state, city, cityPop)

This block groups the result of the previous block by state and calculates the average of the population in each group.

Block Result: Set of Doubles (State, avgCityPop)



Example of a multi-step pipeline

```
db.zipcodes.aggregate( [
  { $group:
    { _id: { gState: "$state", gCity: "$city" },
      cityPop: { $sum: "$pop" } } },
  { $sort: { cityPop: 1 } },
  { $group:
    { _id : "$_id.gState",
      biggestCity: { $last: "$_id.gCity" }, biggestPop: { $last: "$cityPop" },
      smallestCity: { $first: "$_id.gCity" }, smallestPop: { $first: "$cityPop" } } },
  { $project:
    { _id: 0,
      state: "$_id",
      biggestCity: { name: "$biggestCity", pop: "$biggestPop" },
      smallestCity: { name: "$smallestCity", pop: "$smallestPop" } } }
] );
```

For each state, it gets the city with the highest and lowest population. It also shows the population



Example of a multi-step pipeline

```
db.zipcodes.aggregate( [  
  { $group:  
    { _id: { gState: "$state", gCity: "$city" },  
      cityPop: { $sum: "$pop" } } },  
  { $sort: { cityPop: 1 } },  
  { $group:  
    { _id: "$_id.gState",  
      biggestCity: { $last: "$_id.gCity" }, biggestPop: { $last: "$cityPop" },  
      smallestCity: { $first: "$_id.gCity" }, smallestPop: { $first: "$cityPop" } } },  
  { $project:  
    { _id: 0,  
      state: "$_id",  
      biggestCity: { name: "$biggestCity", pop: "$biggestPop" },  
      smallestCity: { name: "$smallestCity", pop: "$smallestPop" } } }  
] );
```

Get the population of each city

Sorts the response of the previous block in crescent order of the population

From the answer of the previous block, it obtains the largest and smallest population and their respective cities

Formats the documents from the previous block response to the final result



Aggregation functions

Examples of functions that can be used in step \$group:

- Same as SQL: \$sum, \$avg,\$min,\$max
- \$push : return a vector with all the values that appear in the group for the chosen field;
- \$addToSet
- \$first: returns the value of an attribute from the first document in an ordered set of documents;
- \$last: returns the value of an attribute from the last document in an ordered set of documents;



Example - \$push and \$addToSet

For each state, it returns a vector with all the cities in the state with repeats and another vector with the cities without repeats

```
db.zipcodes.aggregate(  
  [  
    { $group: { _id: "$state",  
                bagOfCities: { $push: "$city" },  
                setOfCities: { $addToSet: "$city" }  
            }  
    }  
  ]  
);
```



SQL vs Mongodb

| SQL Terms, Functions, and Concepts | MongoDB Aggregation Operators |
|------------------------------------|---|
| WHERE | \$match |
| GROUP BY | \$group |
| HAVING | \$match |
| SELECT | \$project |
| ORDER BY | \$sort |
| LIMIT | \$limit |
| SUM() | \$sum |
| COUNT() | \$sum \$sortByCount |
| join | \$lookup |
| SELECT INTO NEW_TABLE | \$out |
| MERGE INTO TABLE | \$merge (Available starting in MongoDB 4.2) |
| UNION ALL | \$unionWith (Available starting in MongoDB 4.4) |



Examples: SQL vs MongoDB

| SQL Example | MongoDB Example | Description |
|--|--|--|
| <pre>SELECT COUNT(*) AS count FROM orders</pre> | <pre>db.orders.aggregate([{ \$group: { _id: null, count: { \$sum: 1 } } }])</pre> | Count all records from <code>orders</code> |
| <pre>SELECT cust_id, SUM(price) AS total FROM orders GROUP BY cust_id</pre> | <pre>db.orders.aggregate([{ \$group: { _id: "\$cust_id", total: { \$sum: "\$price" } } }])</pre> | For each unique <code>cust_id</code> , sum the <code>price</code> field. |



Examples: SQL vs Mongodb

```
SELECT cust_id,  
       ord_date,  
       SUM(price) AS total  
FROM orders  
GROUP BY cust_id,  
         ord_date
```

```
db.orders.aggregate( [  
  {  
    $group: {  
      _id: {  
        cust_id: "$cust_id",  
        ord_date: { $dateToString:  
          format: "%Y-%m-%d",  
          date: "$ord_date"  
        }  
      },  
      total: { $sum: "$price" }  
    }  
  }  
] )
```

For each
unique
cust_id,
ord_date
grouping,
sum the
price field.
Excludes the
time portion
of the date.



Examples: SQL vs Mongodb

```
SELECT cust_id,  
       ord_date,  
       SUM(price) AS total  
FROM orders  
GROUP BY cust_id,  
         ord_date  
HAVING total > 250
```

```
db.orders.aggregate( [  
  {  
    $group: {  
      _id: {  
        cust_id: "$cust_id",  
        ord_date: { $dateToString  
          format: "%Y-%m-%d",  
          date: "$ord_date"  
        }  
      }  
    },  
    total: { $sum: "$price" }  
  }  
],  
{ $match: { total: { $gt: 250 } } }  
)
```

For each unique `cust_id`, `ord_date` grouping, sum the `price` field and return only where the sum is greater than 250. Excludes the time portion of the date.



Examples: SQL vs MongoDB

```
SELECT cust_id,  
       SUM(price) as total  
FROM orders  
WHERE status = 'A'  
GROUP BY cust_id  
HAVING total > 250
```

```
db.orders.aggregate( [  
  { $match: { status: 'A' } },  
  {  
    $group: {  
      _id: "$cust_id",  
      total: { $sum: "$price" }  
    }  
  },  
  { $match: { total: { $gt: 250 } } }  
)
```

For each unique `cust_id` with status `A`, sum the `price` field and return only where the sum is greater than 250.



Examples: SQL vs MongoDB

```
SELECT COUNT(*)
FROM (SELECT cust_id,
             ord_date
      FROM orders
      GROUP BY cust_id,
              ord_date)
as DerivedTable
```

```
db.orders.aggregate( [
  {
    $group: {
      _id: {
        cust_id: "$cust_id",
        ord_date: { $dateToString
                    format: "%Y-%m-%d",
                    date: "$ord_date"
                  }
      }
    }
  },
  {
    $group: {
      _id: null,
      count: { $sum: 1 }
    }
  }
] )
```

Count the number of distinct `cust_id`, `ord_date` groupings. Excludes the time portion of the date.



Join of collections

- Step \$lookup does an outer join to the left, with an equal condition between the value of an attribute of the documents in the input collection and the value of a field of the documents in the collection to be "joined"

```
{
  $lookup:
  {
    from: <collection to join>,
    localField: <field from the input documents>,
    foreignField: <field from the documents of the "from" collection>,
    as: <output array field>
  }
}
```

```
SELECT *, <output array field>
FROM collection
WHERE <output array field> IN (SELECT *
                                FROM <collection to join>
                                WHERE <foreignField> = <collection.localField>
                                );
```



Examples

```
db.orders.aggregate( [  
  {  
    $lookup:  
    {  
      from: "inventory",  
      localField: "item",  
      foreignField: "sku",  
      as: "inventory_docs"  
    }  
  }  
])
```

```
db.pessoas.aggregate([  
  { $lookup: {  
    from: "compras",  
    localField: "_id",  
    foreignField: "id_comprador",  
    as: "historico_compras" } },  
  { $match: { "historico_compras": { $ne: [] } } }  
]);
```



Join Conditions with Subqueries

```
{
  $lookup:
  {
    from: <joined collection>,
    let: { <var_1>: <expression>, ..., <var_n>: <expression> },
    pipeline: [ <pipeline to run on joined collection> ],
    as: <output array field>
  }
}
```

| Field | Description |
|----------------------|---|
| from | Specifies the collection in the same database to perform the join operation. Starting in MongoDB 5.1, the from collection can be sharded. |
| let | Optional. Specifies variables to use in the pipeline stages. Use the variable expressions to access the fields from the joined collection's documents that are input to the pipeline. |

```
SELECT *, <output array field>
FROM collection
WHERE <output array field> IN (
  SELECT <documents as determined from the pipeline>
  FROM <collection to join>
  WHERE <pipeline>
);
```

To reference variables in the stages pipeline steps, the following syntax is used: "\$\$<variable>"



Example of a pipeline involving join

Query: Get people's data and data on purchases they've made from April 2020

```
db.pessoas.aggregate([
  { $lookup: {
    from: "compras",
    let: { id_pessoa: "$_id" },
    pipeline: [
      { $match: { $expr: {
        $and: [
          { $eq: [ "$id_comprador", "$$id_pessoa" ] },
          { $gte: [ "$data", ISODate("2020-04-01") ] } ] } } },
      { $project: { "id_comprador": 0, "_id": 0 } } ],
    as: "compras_recntes"
  } },
  { $match: { "compras_recntes": { $ne: [] } } }
]);
```

Defines variables for the weighted collection attributes that will be used in the join condition

Join/Selection Condition to tuples of purchases collection:
id_comprador == id_pessoa
Data_compra >= 01/04/2020

Defines the structure of the objects in the Purchases collection that will be included in the join response



Example:

```
db.orders.aggregate( [  
  {  
    $lookup: {  
      from: "restaurants",  
      localField: "restaurant_name",  
      foreignField: "name",  
      let: { orders_drink: "$drink" },  
      pipeline: [ {  
        $match: {  
          $expr: { $in: [ "$$orders_drink", "$beverages" ] }  
        }  
      } ],  
      as: "matches"  
    }  
  }  
])
```

Joins order and restaurant collections by matching localField orders.restaurant_name to ForeignField restaurant.name.

Matching is performed before the pipeline runs.

Performs an array match\$in Between the orders.drink and restaurant.beverages fields. These are accessed using \$\$orders_drink and \$beverages respectively.

```
SELECT *, matches  
FROM orders  
WHERE matches IN (  
  SELECT *  
  FROM restaurants  
  WHERE restaurants.name = orders.restaurant_name  
  AND restaurants.beverages = orders.drink  
);
```



References

MongoDB: The Definitive Guide

Kristina Chodorow

O'Reilly

MongoDB Manual: <http://docs.mongodb.org/manual/>

